

CSCE-771 Project Report: Sentiment Analysis using Large Movie Review dataset

Fawad Kirmani

November 2020

1 Introduction

The goal of this work is to learn and explore different sentiment analysis algorithms. The data used in this work is large movie review dataset prepared by Andrew Maas et al. [1]. For this project, different word vectorization and embedding models are implemented to tokenize the IMDB review dataset. Multiple traditional machine learning models (like Random Forest, Logistic Regression and XGBoost) and deep learning (like CNN and LSTM) classification models have been explored to classify movie reviews into positive and negative classes. Pretrained language models based on transformers (like BERT and DistilBERT) are implemented to learn and achieve better classification scores. State of the art transformer based language models like XLM-RoBERTa and XLNet have also been implemented for this project. The main aim of this project is to learn traditional as well as state of the art methods for sentiment analysis.

In the next section titled (2) *Data*, there is discussion of the data used in this project in more detail. In the section titled (3) *Original work on the dataset*, there is brief discussion of the paper by Andrew Maas et al. [1], which prepared and used this dataset in the paper. In the the section titled (4) *Different approaches for sentiment analysis*, there is a discussion in detail about the implementation and results of methods and algorithms that are explored for sentiment analysis in this project. The work in this project has been divided into three tracks in section (4):

1. Machine Learning models (like Random Forest, Logistic Regression and XGBoost) with TF-IDF vectorization () for classification.

2. Deep Learning models (like CNN and LSTM) with GloVe Word-Embeddings for classification.
3. Classification with Transformers based Pretrained models like BERT, DistilBERT and XLNet.

In the section title (5) Evaluation, this project is evaluated with goal set at the start of the project. In this section, the results are compared with results from Andrew Maas et al. [1] and previous paper on XLNet by Zhilin Yang et al. [7]. In the last section (6) titled Discussion, learning, best results, thoughts and challenges are discussed.

2 Dataset

The data used in this work comes from Andrew Maas et al. [1]. The dataset contains 50,000 movie reviews from IMDB, allowing no more than 30 reviews per movie. The dataset has even number number of positive and negative movie reviews. The dataset contains highly polar movie reviews data. A negative review has a score ≤ 4 out of 10, and a positive review has a score ≥ 7 out of 10. The dataset is evenly divided into training and test sets with 25,000 reviews each. Both the training and test have 12,500 positive and negative reviews. This dataset is widely used to benchmark new work.

For this project, all the positive and negative reviews from both training and test datasets are combined into one dataset. So, there are total 50,000 movie reviews in the dataset. To explore different algorithms for sentiment analysis in this project, positive reviews are labeled as 1 and negative reviews are labeled as 0.

3 Original work on the dataset

In the original work, Andrew Maas et al. [1] presented a model that uses a mix of unsupervised and supervised techniques to learn word vectors capturing semantic term-document information as well as rich sentiment content. Their proposed model leveraged both continuous and multi-dimensional sentiment information as well as non-sentiment annotations. They incorporated the model to utilize the document-level sentiment polarity annotations present in many online documents (e.g. star ratings).

Andrew Maas et al. [1] tested their model on large movie review dataset from IMDB. The model performed best when concatenated with bag of words representation and additional unlabeled dataset. The full model with bag of words and additional unlabeled dataset achieved accuracy of around 89%. The variant of their model which utilized extra unlabeled data during training performed best. They provided the large IMDB movie review dataset publicly to serve as robust benchmark for future work. This project is concentrating on the sentiment analysis part of the Andrew Maas et al. [1] paper where this data from is used.

4 Different approaches for sentiment analysis

4.1 Machine Learning models with TF-IDF vecotization

In the first approach to classify movie reviews, words are vectorized with TF-IDF. Then these vectorized words are used as input to traditional machine learning classification moodels like Logistic Regression, Random Forest and XGBoost.

TF-IDF stands for term frequency–inverse document frequency. The term frequency (TF) is raw count of a term in a document. The inverse document frequency (IDF) is a measure of how much information the word provides, i.e., if it's common or rare across all documents. TF-IDF is the product of two frequencies TF and IDF. Given N documents. For term (word) t in document d

$$Termfrequency(TF) = tf(t, d)$$

$$InversedocumentfrequencyIDF(t) = \log[N/DF(t)] + 1$$

$$TF - IDF(t, d) = TF(t, d) * IDF(t)$$

For TF-IDF transformer, ngrams of 1 to 5 are considered. After fitting ngrams count to tf-idf transformers, total 28,918,046 features are created. The dataset is divided into train and test set of 75% to 25% ratio respectively. All three models i.e., Logistic Regression, Random Forest and XGBoost are trained with GridSearchCV to find best hyperparameters for each model. Table (4.1) shows the precision, recall, f1-score and accuracy on the test dataset with all three machine learning models. Logistic regression model

was overtrained, it is giving accuracy of 90% on training data but only 83% accuracy on test dataset. So, the results of Logistic Regression model are not considered for final comments in this section. The best results are achieved with Random Forest model in this section of the project, where the model achieved precision, recall and f1-score of 84% (macro average). Random Forest model achieved accuracy of 84% as well which is best for this section of the project. The results with XGBoost model were not as good as Random Forest, it achieved precision, recall and f1-score of 82% only. The accuracy of XGBoost model was 82%.

Figures (1) show the normalized confusion matrixes for Logistic Regression, Random Forest and XGBoost classification models. Random Forest model gave the least false negative rate (FNR) of around 20% and false positive rate (FPR) of 12% in this section of the project.

| ML Model | Precision | Recall | F1-score | Accuracy |
|-----------------------------|-----------|--------|----------|----------|
| <i>Logistic Regression*</i> | 83% | 83% | 83% | 83% |
| <i>Random Forest</i> | 84% | 84% | 84% | 84% |
| <i>XGBoost</i> | 82% | 82% | 82% | 82% |

Table 1: Machine Learning models results with TF-IDF vectorization of IMDB movie reviews dataset. The table contains results on test dataset. These are the macro averages of Precision, Recall and F1-score. Logistic Regression model is overtrained, giving training accuracy of 90%.

4.2 Deep Learning with GloVe Word Embeddings

In the second method to classify movie reviews into positive and negative classes, GloVe [2] word embedding is implemented. Then these tokenized word embeddings are used for Deep Learning classification models like Convolutional Neural Network (CNN), Long Short Term Memory (LSTM), Autoencoder LSTM, and Bidirectional LSTM.

GloVe [2] is an unsupervised learning algorithm for obtaining vector representations for words. Training is performed on aggregated global word-word co-occurrence statistics from a corpus, and the resulting representations showcase interesting linear substructures of the word vector space.

The features from GloVe embeddings are trained with different deep learning models. Training and test datasets are divided in 90% to 10% ratio.

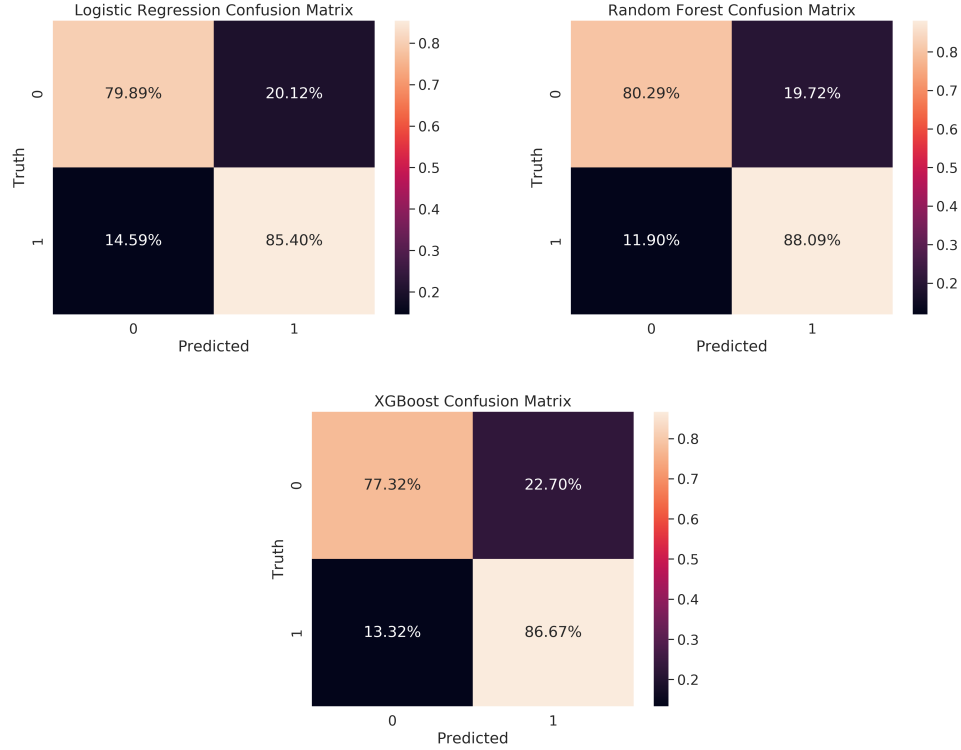


Figure 1: Normalized Confusion matrix for Logistic Regression, Random Forest and XGBoost classification models on TF-IDF vectorized lagre movie review dataset.

These deep learning models are implemented using Keras [9] library available for Python. All the models in this section have been implemented with early stopping function from Keras to check if consecutive epochs are reducing validation dataset loss. If the loss is not decreasing or it is getting worse for few consecutive epochs, model training is stopped to prevent overfitting. Here are the deep learning models implemented with brief description of their implementation -

1. CNN: Figure (2) shows the CNN model summary. The CNN model has one embedding layer. There are two Convolutional layers and four Dense layers in the CNN model. There are 319,601 trainable parameters for CNN model.

2. LSTM: Figure (3) shows the LSTM model summary. The LSTM model has one embedding layer. The LSTM model has two LSTM layers and one output layer. There are 5,305 trainable parameters for LSTM model.
3. Autoencoder LSTM: Figure (4) shows the Autoencoder LSTM model summary. The Autoencoder LSTM model has one embedding layer. The Autoencoder LSTM model has one LSTM layers for encoding and one LSTM layer for decoding. There are 43,583 trainable parameters for Autoencoder LSTM model.
4. Bidirectional LSTM: Figure (5) shows the Bidirectional LSTM model summary. The Bidirectional LSTM model has one embedding layer. The Bidirectional LSTM model has two Bidirectional LSTM layers. There are 9,941 trainable parameters for Bidirectional LSTM model.

CNN gave the best overall results in this section of the project. CNN achieved precision, recall and f1-score 90% (macro average). The accuracy of the CNN model was also the best i.e., 90%. But LSTM model gave the least false negative rate i.e., 7.88%. The least false positive rate was achieved by Bidirectional LSTM of 11.25%. The three models i.e., LSTM, Autoencoder LSTM and Bidirectional LSTM performed similarly in terms of precision, recall, f1-score and accuracy. The detailed results are presented the Table (4.2).

| DL Model | Precision | Recall | F1-score | Accuracy |
|----------------------------|-----------|--------|----------|----------|
| <i>CNN</i> | 90% | 90% | 90% | 90% |
| <i>LSTM</i> | 89% | 89% | 89% | 89% |
| <i>Autoencode LSTM</i> | 89% | 89% | 89% | 89% |
| <i>Bi-directional LSTM</i> | 89% | 89% | 89% | 89% |

Table 2: Deep Learning models results with GloVe word embeddings on IMDB dataset. The classification scores presented in this table are from test dataset. These are the macro averages of Precision, Recall and F1-score.

Figure (6) shows the model accuracy vs epoch, model loss vs epoch, and normalized confusion matrix for the CNN model. Figure (7) shows the model accuracy vs epoch, model loss vs epoch, and normalized confusion matrix graphs for LSTM model. Figure (8) shows the model accuracy

| CNN Model Summary | | |
|---|------------------|----------|
| Model: "sequential_1" | | |
| Layer (type) | Output Shape | Param # |
| embedding_1 (Embedding) | (None, 300, 300) | 29147100 |
| conv1d (Conv1D) | (None, 298, 16) | 14416 |
| max_pooling1d (MaxPooling1D) | (None, 149, 16) | 0 |
| leaky_re_lu (LeakyReLU) | (None, 149, 16) | 0 |
| dropout (Dropout) | (None, 149, 16) | 0 |
| conv1d_1 (Conv1D) | (None, 147, 32) | 1568 |
| max_pooling1d_1 (MaxPooling1D) | (None, 73, 32) | 0 |
| leaky_re_lu_1 (LeakyReLU) | (None, 73, 32) | 0 |
| dropout_1 (Dropout) | (None, 73, 32) | 0 |
| flatten_2 (Flatten) | (None, 2336) | 0 |
| dense_2 (Dense) | (None, 128) | 299136 |
| dropout_2 (Dropout) | (None, 128) | 0 |
| batch_normalization (Batch Normalization) | (None, 128) | 512 |
| leaky_re_lu_2 (LeakyReLU) | (None, 128) | 0 |
| dense_3 (Dense) | (None, 32) | 4128 |
| dropout_3 (Dropout) | (None, 32) | 0 |
| batch_normalization_1 (Batch Normalization) | (None, 32) | 128 |
| leaky_re_lu_3 (LeakyReLU) | (None, 32) | 0 |
| dense_4 (Dense) | (None, 1) | 33 |
| Total params: 29,467,021 | | |
| Trainable params: 319,601 | | |
| Non-trainable params: 29,147,420 | | |
| None | | |

Figure 2: CNN model summary

vs epoch, model loss vs epoch, and normalized confusion matrix graphs for Autoencoder LSTM model. Figure (9) shows the model accuracy vs epoch,

| LSTM Model Summary | | |
|----------------------------------|------------------|----------|
| Model: "sequential_2" | | |
| Layer (type) | Output Shape | Param # |
| embedding_2 (Embedding) | (None, 300, 300) | 29147100 |
| lstm_2 (LSTM) | (None, 300, 4) | 4880 |
| lstm_3 (LSTM) | (None, 8) | 416 |
| dense_5 (Dense) | (None, 1) | 9 |
| Total params: 29,152,405 | | |
| Trainable params: 5,305 | | |
| Non-trainable params: 29,147,100 | | |
| None | | |

Figure 3: LSTM model summary

| Autoencoder LSTM Model Summary | | |
|----------------------------------|------------------|----------|
| Model: "sequential" | | |
| Layer (type) | Output Shape | Param # |
| embedding (Embedding) | (None, 300, 300) | 29147100 |
| lstm (LSTM) | (None, 300, 4) | 4880 |
| flatten (Flatten) | (None, 1200) | 0 |
| repeat_vector (RepeatVector) | (None, 5, 1200) | 0 |
| lstm_1 (LSTM) | (None, 5, 8) | 38688 |
| time_distributed (TimeDistri | (None, 5, 1) | 9 |
| flatten_1 (Flatten) | (None, 5) | 0 |
| dense_1 (Dense) | (None, 1) | 6 |
| Total params: 29,190,683 | | |
| Trainable params: 43,583 | | |
| Non-trainable params: 29,147,100 | | |
| None | | |

Figure 4: Autoencoder LSTM model summary

model loss vs epoch, and normalized confusion matrix graphs for Bidirectional LSTM model.

4.3 Classification with Transformers based Pretrained models

Transformers [3] provides thousands of pretrained models to perform tasks on texts such as classification, information extraction, question answering, summarizing, translation, text generation, etc in 100+ languages. Its aim is to make cutting-edge NLP easier to use for everyone.

| Bidirectional LSTM Model Summary | | |
|----------------------------------|------------------|----------|
| Model: "sequential_3" | | |
| Layer (type) | Output Shape | Param # |
| embedding_3 (Embedding) | (None, 300, 300) | 29147100 |
| bidirectional (Bidirectional) | (None, 300, 8) | 9760 |
| bidirectional_1 (Bidirectional) | (None, 4) | 176 |
| dense_6 (Dense) | (None, 1) | 5 |
| Total params: 29,157,041 | | |
| Trainable params: 9,941 | | |
| Non-trainable params: 29,147,100 | | |
| None | | |

Figure 5: Bidirectional LSTM model summary

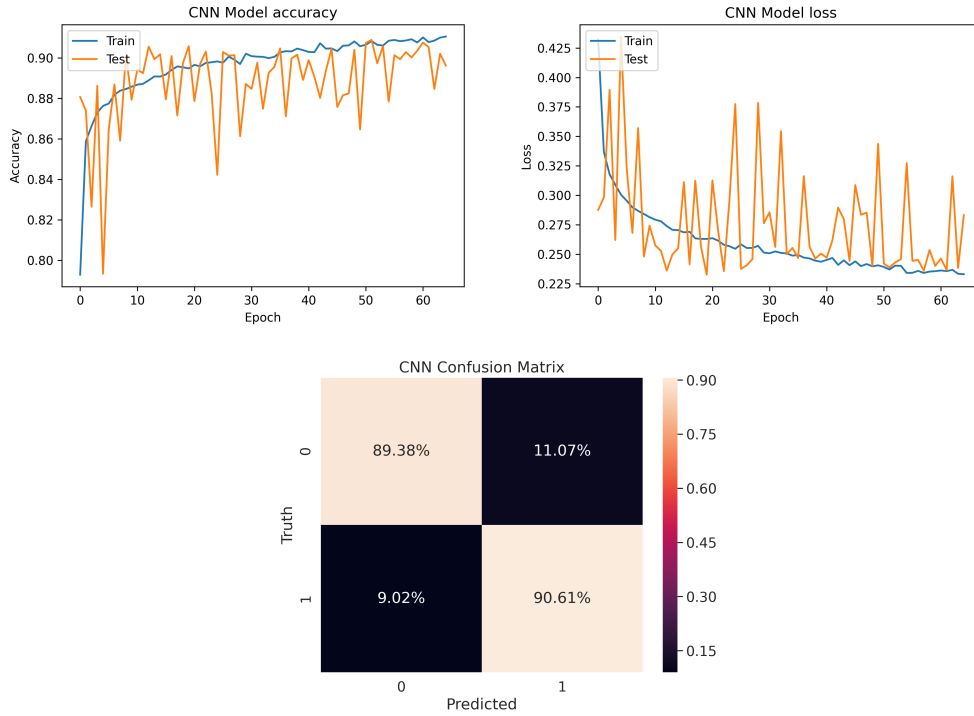


Figure 6: CNN model accuracy vs epoch, CNN model loss vs epoch, and CNN normalized confusion matrix graphs with GloVe word embeddings

This project explored three transformers based pretrained language models for classification of IMDB large movie reviews data:

1. BERT: BERT [5] stands for Bidirectional Encoder Representations




Figure 7: LSTM model accuracy vs epoch, LSTM model loss vs epoch, and LSTM normalized confusion matrix graphs with GloVe word embeddings

from Transformers. BERT [5] is a pre-trained deep bidirectional representations from unlabeled text by jointly conditioning on both left and right context in all layers. As a result, the pre-trained BERT model can be fine-tuned with just one additional output layer to create state-of-the-art models for a wide range of tasks, such as question answering and language inference, without substantial task-specific architecture modifications [5].

2. DistilBERT: DistilBERT [6] is a distilled version of BERT. It is smaller, faster and lighter than BERT. It reduces size of a BERT model by 40%, while retaining 97% of its language understanding capabilities and being 60% faster [6].
3. XLM-RoBETAA: XLM-RoBETAA [1] a transformer-based multilingual masked language model pre-trained on text in 100 languages, which

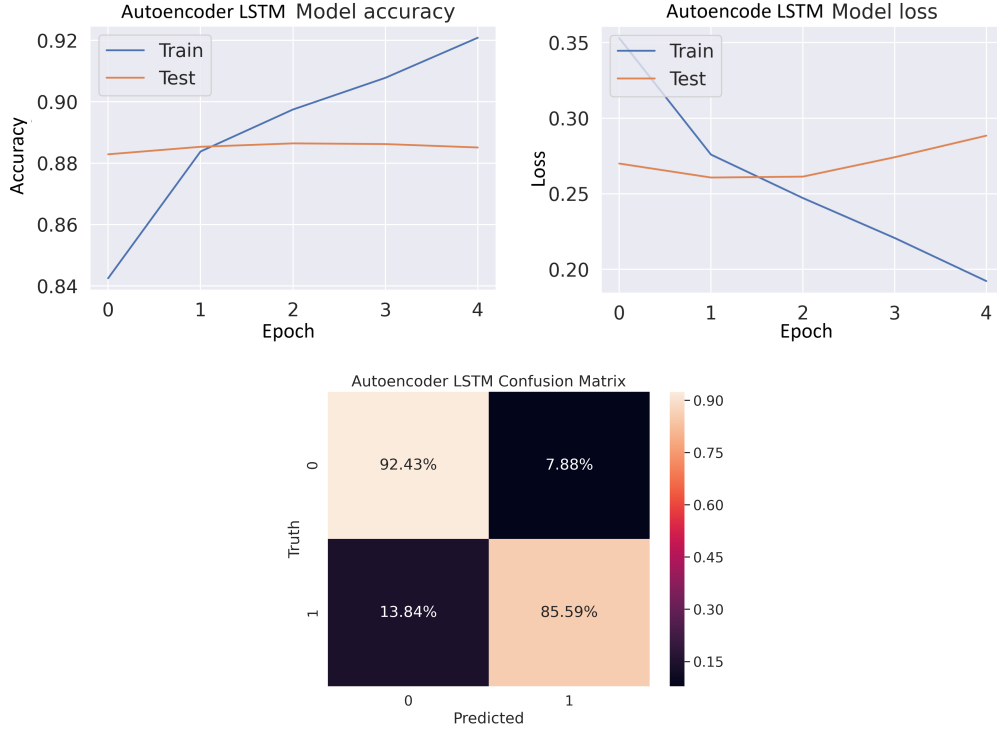


Figure 8: Autoencoder LSTM model accuracy vs epoch, Autoencoder LSTM model loss vs epoch, and Autoencoder LSTM normalized confusion matrix graphs with GloVe word embeddings

obtains state-of-the-art performance on cross-lingual classification, sequence labeling and question answering. XLM-RoBERTa (XLM-R) outperforms mBERT on cross-lingual classification by up to 23% accuracy on low-resource languages [10].

4. XLNet: XLNet [7] is a generalized autoregressive pretraining method that (1) enables learning bidirectional contexts by maximizing the expected likelihood over all permutations of the factorization order and (2) overcomes the limitations of BERT thanks to its autoregressive formulation [7].

This project used ktrain [8] library to train transformers based pretrained models. ktrain is low-code library for augmented machine learning. It is a wrapper to TensorFlow and many other libraries (e.g., transformers, scikit-

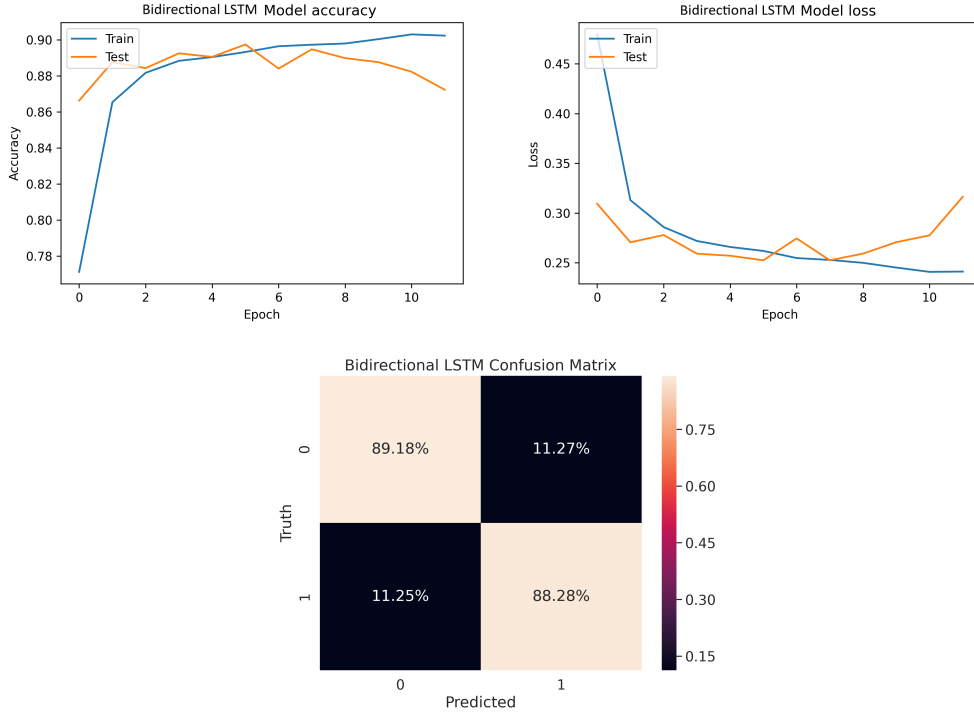


Figure 9: Bidirectional LSTM model accuracy vs epoch, Bidirectional LSTM model loss vs epoch, and Bidirectional LSTM normalized confusion matrix graphs with GloVe word embeddings

learn, stellargraph). It is one of the latest state of the art library to train machine learning models.

To fine tune BERT, DistilBERT, XLM and XLNet classifiers, learning rate is tuned for each of them using in-built learning rate finder `lr_find()` function in `ktrain` [8]. Figure (10) shows the graph generated for loss vs learning_rate (log scale) to find best learning rate for BERT model. BERT [5] paper uses learning rate of $5e-5$, $4e-5$, $3e-5$ and $2e-5$. I choose learning rate of $2e-5$ as in the figure (10), we can see the loss is going down at learning rate of $2e-5$. BERT model is trained with max length of 256 and 3 epochs. BERT model achieved precision, recall, f1-score and accuracy scores of 93%.

To explore more on learning late and `lr_find()`, DistilBERT model was first trained with learning rate of $2e-4$ and 4 epochs after observing the plot in figure (10). DistilBERT, didn't achieved good results with for max length

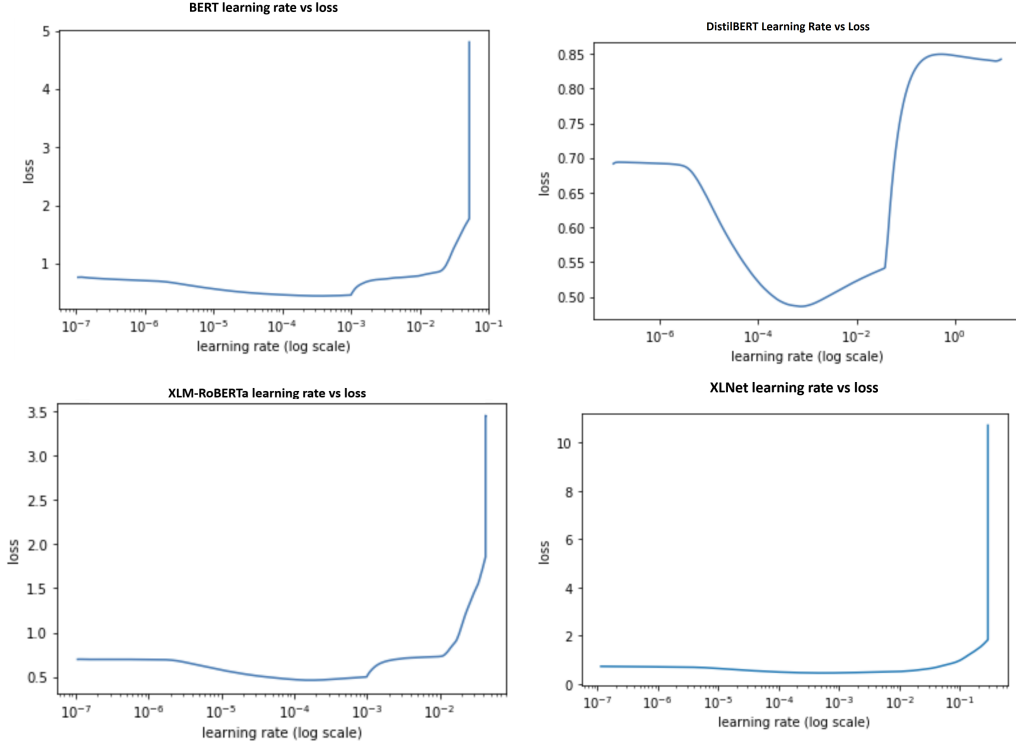


Figure 10: BERT, DistilBERT, XLM-Roberta and XLNet learning rate vs loss graphs with 1, 2, 2 and 1 epochs respectively.

of 500 and learning rate of $2e-4$ and 4 epochs. The model got overtrained and achieved accuracy of 90% on test data while 97% on training data. Then, DistilBERT model was trained on max length of 500 with learning rate of $2e-5$ and 3 epochs, the model achieved pretty good precision, recall and f1-score scores of 94%. DistilBERT also achieved accuracy score of 94%.

XLM-RoBERTa model was first trained with max length of 256 and learning rate of $2e-4$ and 4 epochs. XLM-RoBERTa model got overtrained with 4 epoch, on training data it gave accuracy of 96% but on testing data it gave accuracy of only 90%. Then, XLM-RoBERTa model was trained on max length of 512 with learning rate of $2e-5$ and 2 epochs, the model achieved pretty good precision, recall and f1-score scores of 94%. XLM-RoBERTa also achieved accuracy score of 94%.

For XLNet, it was taking very long time to explore training on different learning rates, so to complete project on time, XLNet model with max length

of 512 was trained with learning rate of $2e-5$ with 3 epochs similar to DistilBERT. The best results are achieved with XLNet. XLNet model achieved precision, recall, f1-score and accuracy scores of 95%. Table (4.3) shows the scores achieved by all the three (BERT, DistilBERT and XLNet) pretrained transformers based models implemented for this project.

Table (4.3) shows the results for BERT, DistilBERT, XLM-RoBERTa and XLNet. Classification with BERT achieved precision, recall, f1-score and accuracy of 93%. Classification with DistilBERT achieved precision, recall, f1-score and accuracy of 94%. Classification with XLM-RoBERTa achieved precision, recall, f1-score and accuracy of 94%. XLNet achieved the best classification scores. XLNet achieved precision, recall, f1-score and accuracy of 95%. XLNet model achieves lowest FPR of 4.59% and lowest FNR of 5.19% among all the models implemented for this project.

| Transformers | Precision | Recall | F1-score | Accuracy |
|---------------------|------------------|---------------|-----------------|-----------------|
| <i>Bert</i> | 93% | 93% | 93% | 93% |
| <i>Distilbert</i> | 94% | 94% | 94% | 94% |
| <i>XLM-Roberta</i> | 94% | 94% | 94% | 94% |
| <i>XLNet</i> | 95% | 95% | 95% | 95% |

Table 3: Sentiment classification of movie reviews dataset with Transformer based pretrained language models. These are the macro averages of Precision, Recall and F1-score.

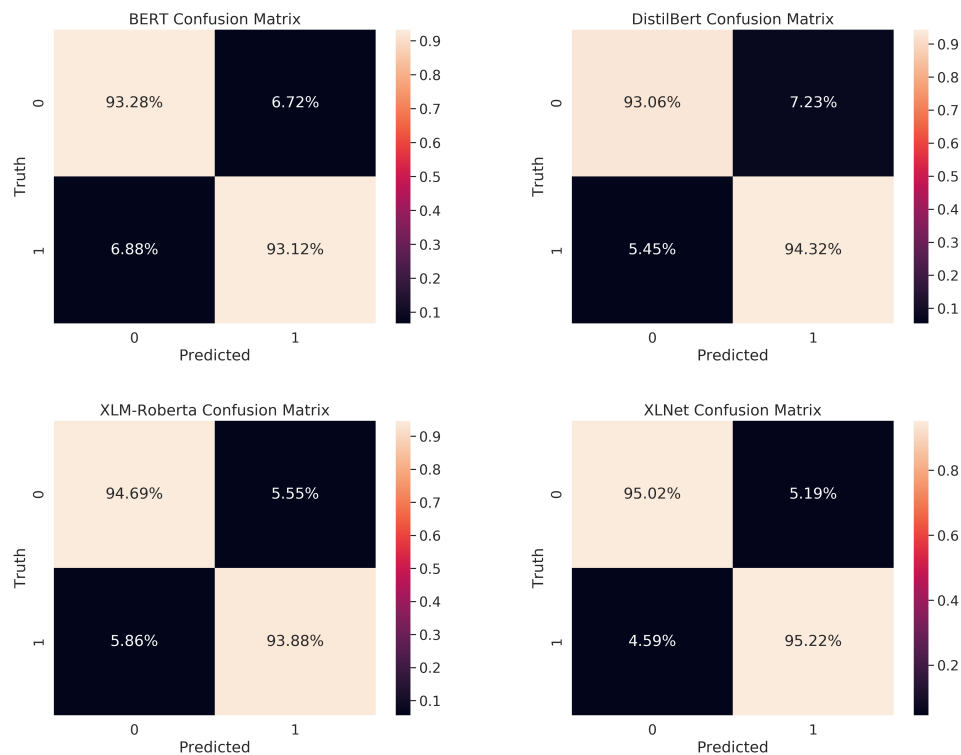


Figure 11: BERT, DistilBERT, XLM-Roberta and XLNet normalized confusion matrix on IMDB large movie reviews dataset

5 Evaluation

The main goal of the project was to learn and implement different natural language processing techniques (NLP) for sentiment analysis. The project mostly accomplishes its goal. The only thing that could have been better if the project could further explore another word embedding method and learn more about hyperparameter tuning of pretrained language models. But given the time allotted for the course, almost all the objectives of the project are accomplished.

The other goal of the project was to measure the performance of the results of the sentiment classification model implemented in the project to Andrew Maas et al. [1]. In that sense -

1. The machine learning models trained on TF-IDF vectorized dataset

didn't perform well at all. The best model among them i.e., Random Forest model gave accuracy of only 84%. It was lower than 5% points than Andrew Maas et al. [1] accuracy score of 89%.

2. The deep learning models trained on GloVe word embeddings datasets did perform as well. The best model among them i.e., CNN gave slightly better accuracy than Andrew Maas et al. [1] accuracy score of 89%. Other models i.e., LSTM, Autoencode LSTM and Bi-directional LSTM gave same accuracy score of 89% as Andrew Maas et al.
3. The transformers based pretrained models performed best. They outperformed Andrew Maas et al. [1] accuracy score by 5-6% points. XLNet model performed best and achieved accuracy of 95% compared to Andrew Maas et al. score of 89%. Other transformer based pretrained language models also performed better than Andrew Maas et al.

The transformer based language based models can be further improved by tuning hyperparameters with more GPU resources or time or both. For example, XLNet [7] model on the same movie reviews data achieved 96% accuracy. So, with little more hyperparameter tuning similar performance can be achieved.

6 Discussion

In this project, different methods and algorithms were implemented for sentiment analysis on large movie reviews dataset [1]. From TF-IDF tokenizer, GloVe word embeddings to state of the art pretrained language model like XLNet. Machine learning models implemented on TF-IDF tokenizer didn't performed well, giving best accuracy of 84% with Random Forest classification. Deep Learning models implemented with Glove word embeddings performed similar to Andrew Maas et al. [1], giving best accuracy score of 90% with CNN classification. Transformer based learning model performed best, achieving best accuracy 95% with XLNet classification model.

The main learning's from the project was to learn text tokenizing, word embeddings and transformers based pretrained models. The main challenge was to tune hyperparameters for transformers. To tune one hyperparameter, it used to take lot of hours, sometime days. So, only the parameter which was explored in some ways was learning rate for different transformers based

models. The performance of transformer based pretrained language based models could have been further improved with some hyperparameter tuning as shown in Sun et al. [11]. That is one of the future tasks to explore.

This project shows how the text based sentiment analysis has come a long way during last few years through these algorithms. The project shows that transformer based pretrained language models are performing significantly better than traditional machine learning and deep learning models.

References

- [1] Maas, Andrew & Daly, Raymond & Pham, Peter & Huang, Dan & Ng, Andrew & Potts, Christopher. (2011). Learning Word Vectors for Sentiment Analysis. 142-150.
- [2] Pennington, Jeffrey & Socher, Richard & Manning, Christopher. (2014). GloVe: Global Vectors for Word Representation. EMNLP. 14. 1532-1543. 10.3115/v1/D14-1162.
- [3] Vaswani, Ashish & Shazeer, Noam & Parmar, Niki & Uszkoreit, Jakob & Jones, Llion & Gomez, Aidan & Kaiser, Lukasz & Polosukhin, Illia. (2017). Attention Is All You Need.
- [4] F. Pedregosa, et al. (2011) Scikit-learn: Machine Learning in Python. J. Mach. Learn. Res. 12 (November 2011), 2825-2830.
- [5] Devlin, Jacob & Chang, Ming-Wei & Lee, Kenton & Toutanova, Kristina. (2018). BERT: Pre-training of Deep Bidirectional Transformers for Language Understanding.
- [6] Sanh, Victor & Debut, Lysandre & Chaumond, Julien & Wolf, Thomas. (2019). DistilBERT, a distilled version of BERT: smaller, faster, cheaper and lighter.
- [7] Yang, Zhilin & Dai, Zihang & Yang, Yiming & Carbonell, Jaime & Salakhutdinov, Ruslan & Le, Quoc. (2019). XLNet: Generalized Autoregressive Pretraining for Language Understanding.
- [8] Maiya, Arun. (2020). ktrain: A Low-Code Library for Augmented Machine Learning.

- [9] Chollet, F., & others. (2015). Keras. GitHub. Retrieved from <https://github.com/fchollet/keras>
- [10] Conneau, Alexis & Khandelwal, Kartikay & Goyal, Naman & Chaudhary, Vishrav & Wenzek, Guillaume & Guzman, Francisco & Grave, Edouard & Ott, Myle & Zettlemoyer, Luke & Stoyanov, Veselin. (2020). Unsupervised Cross-lingual Representation Learning at Scale. 8440-8451. 10.18653/v1/2020.acl-main.747.
- [11] Chi, Sun & Qiu, Xipeng & Xu, Yige & Huang, Xuanjing. (2019). How to Fine-Tune BERT for Text Classification?.